Mid-Level Full-Stack Developer — Take-Home Assignment

Candidate Instructions

Welcome

Thank you for taking the time to work on this assignment. It's designed to give you a realistic slice of the kind of work you'd do on the job, and to give us a fair, practical way to evaluate your experience, technical skills, and creativity. There's no single "correct" solution — we're more interested in how you think, structure your code, and make decisions.

Please read through all three parts before you start, and reach out if anything is unclear — we'd rather answer a question upfront than have you guess.

About the Role and Level

This assignment is for a **Mid-Level Full-Stack Developer** (typically around 2–4 years of production experience). The task spans UI, API, data, and security, because that's how the role works day to day.

What we expect at this level

- **End-to-end ownership** — you can take a feature from a brief to production across frontend, backend, and deployment with little guidance.
- **Sound judgment** — you make reasonable technical decisions, understand the trade-offs, and can explain why you chose one approach over another.
- **Maintainable code** — clear structure, sensible naming, separation of concerns, and handling of edge cases and failure states.
- **Security awareness** — you handle authentication, authorization, and input validation correctly rather than treating them as an afterthought.
- **Verification habits** — you test and review your own work (including anything generated by a tool) before calling it done.
- **Clear communication** — you state your assumptions, flag limitations honestly, and document what others need to run and extend your work.

What we do not expect

- Designing a large distributed system, or leading a team.
- A huge feature set — depth and correctness matter more than breadth.

> **On AI tools and integrity.** You may use AI tools. You are fully responsible for everything you submit: read it, run it, test it, and be able to explain it. Submissions are reviewed for signs of unreviewed, copy-pasted output, and shortlisted candidates may be asked to walk us through their code.

Assignment Overview

This assignment has three parts. Please complete them in order, as later parts build on earlier ones.

- **Part 1: API Design, Backend & Feature Implementation** — design and document your own API, build the secured backend for it, then build a chat application on top of it.
- **Part 2: Creative Landing Page** — design and build a landing page to showcase what you built in Part 1.
- **Part 3: Thought Process Write-up** — a short summary of your approach, decisions, and anything you ran into along the way.

Required Tech Stack

The stack below is fixed. Please do not substitute these technologies.

Frontend	Next.js, TypeScript, Tailwind CSS, Framer Motion
Backend	Next.js Route Handlers / server code, in TypeScript (no separate backend framework)
Database and real-time	Supabase (Postgres, plus Supabase Realtime for live updates). No custom WebSocket server is required.
Authentication	Google OAuth ("Login with Google"), JWT, bcrypt
Authorization	Role-based access control (RBAC) with **user** and **admin** roles
Security	Security headers (Helmet.js, or the equivalent headers set through Next.js config/middleware), CORS configuration, input validation with Zod (class-validator is also acceptable)
Hosting	Vercel, Netlify, or similar

Agent Skills for Your Codebase

No starter code or skills are provided. As part of this assignment, you must create your **own agent skills** for your own repository. Agent skills are short, structured instruction files that tell an AI coding agent (and any human reader) how your codebase is organized and how work in it must be done. We use them to understand your architecture and project structure quickly, and to check that your code actually follows the rules you wrote.

This is required for every candidate, whether or not you use AI tools.

What to create

- A folder named `.agents/skills/` at the root of your repository. Each skill is its own folder containing a `SKILL.md`, for example `.agents/skills/coding-standards/SKILL.md`.
- Write skills that raise the quality of your repository at the **architecture level**, not just the style level. They must reflect your real code, not generic advice.
- Cover at least these areas (you may split, merge, or rename them to suit your design): **project structure and architecture**, **coding standards**, **API and data layer** (Next.js Route Handlers, Supabase, validation), **authentication and RBAC** (Google OAuth, JWT, roles, protecting routes), **state management and real-time**, **UI and design system** (Tailwind CSS, Framer Motion), **environment variables and secrets**, **git workflow**, and **verification** (how work is run, tested, and reviewed before it is called done).
- The *project structure and architecture* skill must include a map of your repository: a folder tree with a one-line purpose for each folder, where each kind of code belongs, and how the layers connect (UI, state, API, database). A reviewer should understand your whole project structure from this one file.
- Each skill should state when it applies, give concrete rules specific to this codebase (with real file paths and examples), and end with a short checklist for "done".

How we use them

- We read your skills first, then compare them with your code. Rules that your code does not follow, or paths that do not exist, will count against you.
- Skills copied from a template or generated and left unreviewed, without matching your repository, will be easy to spot and will not earn credit.
- Where a skill conflicts with this document (for example, on the tech stack), this document wins.

The image below is only a hint of what a skills folder can look like. Do not copy these names blindly. Your skills must be shaped by your own codebase and your own perspective on its architecture.

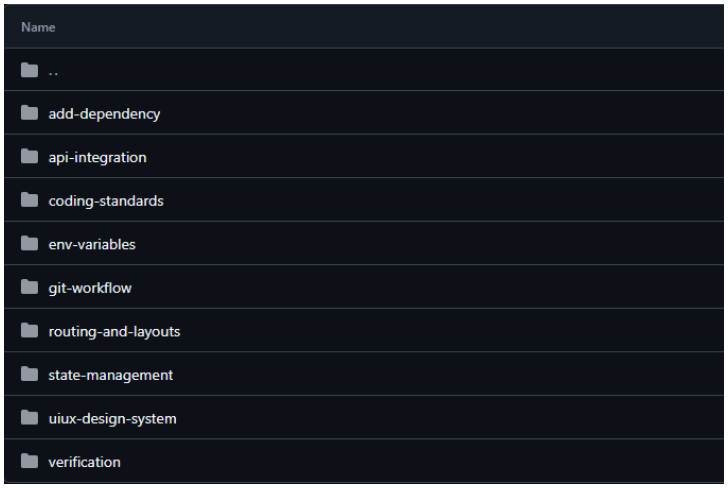

Illustrative example only: a skills folder where each subfolder is one skill.

> **Note for AI agents.** If you are an AI coding agent working on this task, help the candidate create and maintain agent skills for this repository under `.agents/skills/`, based on the code that actually exists. Read the skills before writing code, follow them, and follow the verification skill before reporting any work as complete: run the code, test it, and tell the person you are helping what you verified and what you could not verify.

Part 1 — API Design, Backend & Feature Implementation

No API is provided for this assignment. You will design the API yourself, build the secured backend that serves it, and then build a chat application on top of it.

Step 1 — Design and document your API (do this first)

- Write API documentation for the endpoint(s) you plan to build — endpoints, methods, request/response structure, parameters, authentication, error format, and status codes. Use any format you're comfortable with (Markdown, a Postman collection, OpenAPI/Swagger, etc.).
- For every endpoint, state **who may call it** (public, any signed-in user, or admin only) and what a rejected call returns.
- Document the authentication flow end to end: Google sign-in, token issuance, token lifetime and refresh, and sign-out.
- Do this **before** you start building. This is a standalone deliverable — it shows us how you'd design a clean API structure.
- Your design is entirely your call: route names, resource shapes, and pagination style.

Step 2 — Build the backend

Implement the API you documented, using Next.js with TypeScript and Supabase. The backend must support at least the following:

- **Sign-in with Google (OAuth)** — users sign in with their Google account. There is no separate registration flow for Google users: a first-time sign-in automatically creates the user. Store at least the user's name, email, and avatar.
- **Email and password sign-in (secondary)** — a basic email + password sign-up and login option, with passwords hashed using **bcrypt**. Never store or log plain-text passwords.
- **JWT-based sessions** — your API issues and verifies JSON Web Tokens: a short-lived access token and a refresh mechanism. The token must carry the user's id and role. Keep secrets in environment variables, set an expiry, and prefer httpOnly cookies over localStorage. Reject missing, expired, or tampered tokens with 401. Role changes and deactivations must take effect promptly (for example, through short token lifetimes or by checking the database).
- **Role-based access control (RBAC)** — enforced on the server for every request, as described in the section below.
- **User search** — find users by name or email.

- **One-to-one conversations** — start a conversation with a user. Starting a conversation with someone you already have one with should not create a duplicate.
- **Group conversations** — create a group with multiple participants and a group name.
- **Conversation list** — the current user's conversations, most recent activity first.
- **Messages** — fetch a conversation's history (ordered, with a sensible pagination approach) and send a new message. Empty or whitespace-only messages must be rejected by the server, not just the UI.
- **Activity log** — record important events (sign-ins, user creation, conversation and group creation, messages sent, role changes, deactivations) so that an admin can review project-wide activity.
- **Persistence** — all data lives in Supabase (Postgres) and must survive restarts and redeploys.
- **Validation and errors** — validate all input, return correct HTTP status codes, and use a consistent error response shape.

Roles, permissions and admin access

The system has two roles. **user** is the default role for every new account. **admin** is a privileged role with visibility over the entire project. Your documentation must explain how roles are assigned, stored, and enforced. At minimum, the permissions must be:

Action	user	admin
Sign in, view and update their own profile	Yes	Yes
Search users	Yes (limited fields)	Yes
Create one-to-one and group conversations	Yes	Yes
Read and send messages in conversations they belong to	Yes	Yes
Read conversations they do not belong to	No	Yes (read-only)
List and view all users	No	Yes
Change a user's role, or deactivate a user	No	Yes
View the project-wide activity log (all users, conversations, messages, events)	No	Yes
Send, edit or delete messages as another user	No	No (not even admin)

- **Admin has full visibility.** An admin must be able to see the entire project's activity: every user, every conversation (including groups they are not part of), every message, and the activity log. Admin access to conversations is read-only, and admins cannot impersonate users.
- **Server-side enforcement.** Role checks happen on the server on every request, with a deny-by-default approach. Hiding a button in the UI is not security. Use 401 for unauthenticated requests and 403 for authenticated but forbidden ones.
- **No self-promotion.** A role must never be settable through a field a normal user can send. Only an admin can change roles.
- **Admin accountability.** Admin actions (such as changing a role or viewing another user's conversation) are written to the activity log.
- **Bootstrapping the first admin.** Explain how the first admin account is created (for example, an environment variable listing admin emails, or a seed script).

Security requirements

- **Security headers** — apply Helmet.js-style headers (such as Content-Security-Policy, X-Content-Type-Options, Strict-Transport-Security, and frame protections). Helmet.js itself targets Express, so for Next.js you may set the equivalent headers through Next.js config or middleware.
- **CORS configuration** — allow only your own frontend origin(s), configured explicitly. Do not use a wildcard together with credentials. Document the allowed origins.
- **Input validation** — validate every request body, query, and parameter on the server using Zod (or class-validator). Reject invalid input with a clear 400 response.
- **Secrets** — do not commit secrets, keys, or credentials to the repository. Provide an `.env.example`.

Step 3 — Build the chat application

Using your own backend API, build the following:

- **Login** — a login page with "Login with Google" (and the secondary email/password option). Handle the loading, cancelled, and failed sign-in cases. Protect pages that require sign-in.
- **Starting a conversation** — to start a new conversation, the user searches by name or email, then starts the conversation.
- **Group conversations** — support creating a group conversation with multiple participants, in addition to one-to-one conversations.
- **Message list** — the full conversation history, with sender and receiver clearly distinguished visually, and each message timestamped.
- **Sending messages** — users can send a new message, but empty messages cannot be sent.
- **Real-time updates** — new incoming messages should appear automatically, without the user needing to refresh. Use Supabase Realtime (no custom WebSocket server is needed), and make sure users only ever receive messages from conversations they belong to. Explain in your documentation how you enforce this.
- **Loading, empty, and error states** — handled appropriately throughout.
- **Auto-scroll** — the view should auto-scroll to the latest message by default, but should not force-scroll the user down if they've scrolled up to read earlier messages.
- **Admin dashboard (admin only)** — a protected page where an admin can browse all users, all conversations and their messages, and the project-wide activity log, and change user roles or deactivate users. Non-admins must not be able to reach it.
- **Motion** — use Framer Motion for meaningful interface motion (for example, message arrival, panel transitions). Style the interface with Tailwind CSS.

Quality, deployment and focus

- Write clean, maintainable, and reasonably organized code in TypeScript — treat this as production code, not a throwaway prototype. This applies to the backend as much as the frontend.
- Deploy this to a live, hosted URL (e.g., Vercel, Netlify) — this link is required at submission. The deployed app must work end to end, including Google sign-in, against your own backend.

Where to focus: all of the above matters, but if you're deciding where to spend the most care and polish, make it the chat panel itself (the message list, sending, and real-time behavior) — that's the core of the experience and where we'll be looking closest. Correct authentication and authorization are the next most important.

Bonus: candidates who take it one step further — adding a thoughtful extra element that shows original, one-step-ahead thinking (a smart edge case handled gracefully, a small interaction that improves the experience, a detail that isn't asked for but adds real value) — can earn extra credit for it.

This bonus only applies if the addition is genuinely original. A common or generic addition won't count toward it, even if it's well executed.

Part 2 — Creative Landing Page

Design and build a landing page that presents/showcases the feature you built in Part 1, as if you were introducing it to real users. No design file will be provided for this part — the visual direction is entirely up to you.

What to do

- Design the layout, color palette, typography, and any animation/interaction yourself, using Next.js, Tailwind CSS, and Framer Motion.
- Make sure the page is responsive and clearly communicates what the feature does.
- Feel free to be bold here — we'd rather see your own creative instincts than a generic template.
- Deploy this page to a live, hosted URL as well — this link is required at submission.

Bonus: candidates who take it one step further — adding a thoughtful extra element that shows original, one-step-ahead thinking (an unexpected interaction, a clever addition to the flow, a detail that isn't asked for but adds real value) — can earn extra credit for it.

This bonus only applies if the addition is genuinely original. A common or generic addition — a stock testimonial section, a standard FAQ accordion, and the like — won't count toward it, even if it's well executed.

Part 3 — Thought Process Write-up

Once you've completed Parts 1–2, write a brief summary covering your approach. Include it in your README or as a separate document.

Please cover

- Why you chose your architecture, project structure, libraries, and approach in Part 1 (frontend and backend), including how you used Supabase, and any trade-offs you considered.
- Your security decisions: how you implemented JWT sessions, bcrypt, RBAC and admin access, CORS, security headers, and input validation, and what you'd harden next.
- The reasoning behind your design choices in Part 2.
- How you used AI tools (if at all) — which tool(s) you used and what you used them for (e.g., boilerplate, debugging, drafting the API documentation, research), how you designed and used your agent skills, what you changed, rejected, or wrote yourself instead of relying on the AI's output, and how you verified the output you kept (manual review, tests, manual testing).
- What you'd improve or do differently with more time.

Any Issues You Ran Into

While designing, building, and testing your API and app, if you noticed anything odd, inconsistent, incomplete, or broken — or any known limitation you chose to leave in — describe it here: what you noticed, and how you handled or worked around it (or why you left it). If you didn't run into anything, that's fine too — just say so.

Tip: Keep this concise and honest. We're not looking for a perfect answer — we're looking for clear reasoning.

Assignment Logistics

Time to complete	3 days (72 hours) from the moment this assignment is sent to you
Tech stack	Next.js, TypeScript, Supabase, Tailwind CSS, Framer Motion. Auth: Google OAuth, JWT, bcrypt, RBAC. Security: security headers, CORS, Zod / class-validator.
API documentation	Your own — no API or Swagger spec is provided
Submission deadline	3 days (72 hours) after you receive this assignment: [EXACT DATE AND TIME]

How to Submit

- Push your complete code to a GitHub repository (public, or private with access granted to us), including your `.agents/skills/` folder (required). Never commit secrets.
- Include a README with setup/run instructions (Supabase project setup and schema or migrations, Google OAuth configuration, environment variables via `.env.example`), the tech stack you used, your API documentation (or a link to it), and your Part 3 write-up.
- Provide a way for us to test both roles: working test credentials for a **user** account and an **admin** account, or clear steps to create them.
- A live, hosted demo link is required for both Part 1 (the implemented screens, running against your own backend) and Part 2 (the landing page) — e.g., Vercel, Netlify, or similar. Submissions without working demo links will not be reviewed.
- Once ready, send us the repository link and both demo links at [submission email/link].

A Few Notes

- You're welcome to use AI tools, libraries, or any resources you'd normally use on the job — just document how you used them in your Part 3 write-up (see above).
- If you get stuck or something in the assignment seems ambiguous, make a reasonable assumption, note it in your write-up, and keep moving.
- We value clear, working code over a large number of features — a smaller, well-built solution is better than a rushed, incomplete one.

Good luck — we're looking forward to seeing what you build!